

OpenRISC Instruction Set Architecture Extensions

Master Thesis

Zurich, May 2015

Andreas Traber

Advisors


Michael Gautschi

Antonio Pullini

Prof. Dr. Luca Benini

ETH zürich

Integrated Systems Laboratory




Outline

- 1 Motivation
- 2 Vectorial ALU
- 3 Improved MAC Unit
- 4 Bit Counting
- 5 Results
- 6 Conclusion

Motivation

Ultra-Low-Power Computing

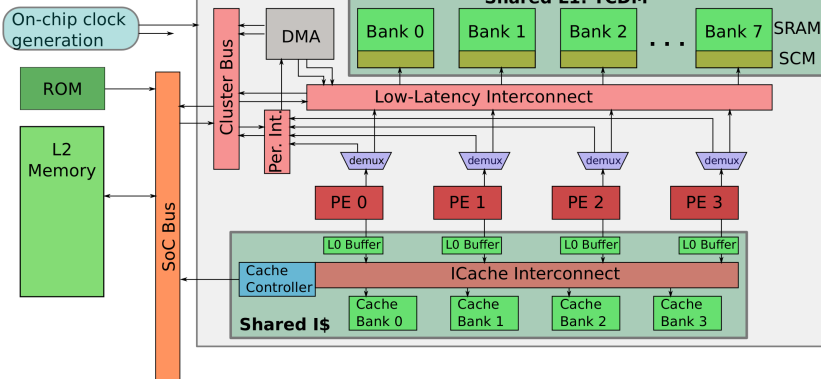
- Today's embedded devices need a vast amount of computing power
Internet of things, wearable devices, sensors, smartphones, ...
- Need to be energy efficient
Operate digital circuits near the threshold voltage
- Use parallelization to recover performance
Turn off unused processing elements



PULP

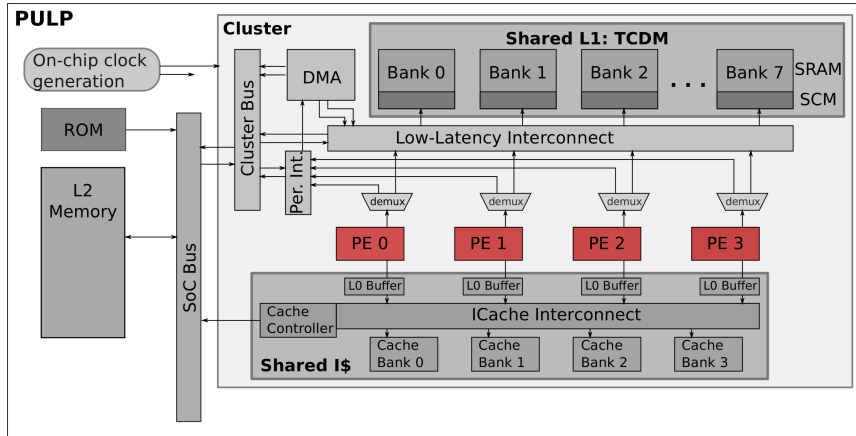
Parallel Ultra-Low power Processor

PULP



PULP

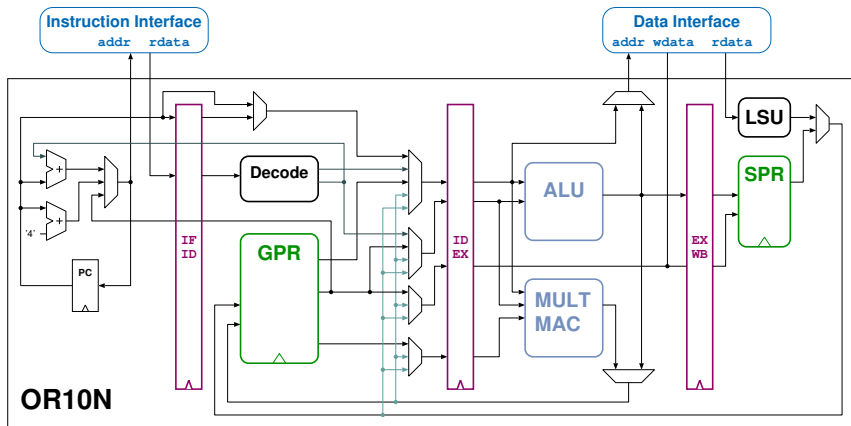
Parallel Ultra-Low power Processor



- My work is focused on the PE

PE: OR10N

In-order 4-stage OpenRISC CPU



Goals

- Higher energy efficiency

Goals

- Higher energy efficiency
- Doing more computations per cycle \Rightarrow ISA extensions
 - Hardware loops
 - Auto-incrementing load/stores
 - Vectorial instructions
 - Improved MAC
 - Bit counting instructions

Goals

- Higher energy efficiency
- Doing more computations per cycle \Rightarrow ISA extensions
 - Hardware loops
 - Auto-incrementing load/stores
 - Vectorial instructions
 - Improved MAC
 - Bit counting instructions
- No increase in clock cycle time

Goals

- Higher energy efficiency
- Doing more computations per cycle \Rightarrow ISA extensions
 - Hardware loops
 - Auto-incrementing load/stores
 - Vectorial instructions
 - Improved MAC
 - Bit counting instructions
- No increase in clock cycle time
- Extensions supported by compiler
 - New instructions generated automatically

Vectorial ALU


Use subword parallelism

- Re-use existing ALU and data path
- Segment the data path dynamically into 2/4 parts
- No separate register file needed

Vectorial ALU

Use subword parallelism

- Re-use existing ALU and data path
- Segment the data path dynamically into 2/4 parts
- No separate register file needed
- Unaligned memory access



Vectorial ALU

Encoding

31							0	
001010	DDDDDD	AAAAA	BBBBB	---	00	sub-opcode	0	lv.inst.h rD, rA, rB
001010	DDDDDD	AAAAA	BBBBB	---	01	sub-opcode	0	lv.inst.h.sc rD, rA, rB
001010	DDDDDD	AAAAA	IIIII	III	10	sub-opcode	0	lv.inst.h.sci rD, rA, I
001010	DDDDDD	AAAAA	BBBBB	---	00	sub-opcode	1	lv.inst.b rD, rA, rB
001010	DDDDDD	AAAAA	BBBBB	---	01	sub-opcode	1	lv.inst.b.sc rD, rA, rB
001010	DDDDDD	AAAAA	IIIII	III	10	sub-opcode	1	lv.inst.b.sci rD, rA, I

■ Sub-opcodes

- | | | | | |
|--------------|--------------|-------|-------|------------------------------|
| ■ ADD | ■ MIN | ■ SRL | ■ OR | ■ INS
■ EXT |
| ■ SUB | ■ MAX | ■ SRA | ■ AND | |
| ■ AVG | ■ ABS | ■ SLL | ■ XOR | |

Vectorial ALU

Comparisons

■ `lv.cmp_*`

Vectorial comparison

■ `lv.all_*`

Set flag if all comparisons evaluate to true

■ `lv.any_*`

Set flag if any comparison evaluates to true

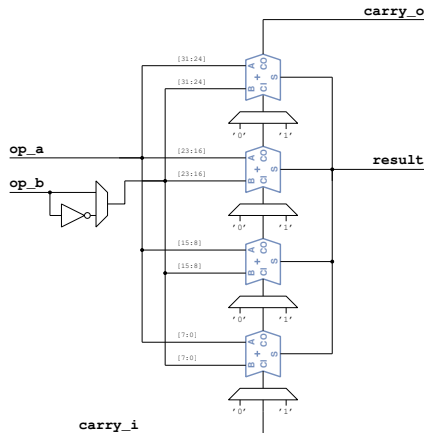
`lv.cmp_ge.b rD, rA, rB`

rA	11011110	10101101	10111110	10101111
	>	>	>	>
rB	10110001	01101011	00000000	10110101
	=	=	=	=
rD	11111111	00000000	00000000	11111111

Vectorial ALU



Simple Example: Vectorial ADD/SUB

- For SUB: Invert operand b
- Four slices chained together
- Suppress carries where needed
0 for ADD
1 for SUB



Improved MAC Unit

- Use only 32-bit multiplication result
- Put accumulator in GPR
 - Multiple accumulators in GPR
 - Instead of one in SPR
- Vectorial MAC
 - Offers same parallelism as ALU
- 16-bit subword selection
 - Accelerates 64-bit multiplication
- Single-cycle execution
 - No pipeline stalls
 - Two-cycles with old MAC



Bit Counting


- Population Count: Number of bits set to 1
- Count leading bits
- Find first/last 1 in a word
- Not automatically generated by compiler
Need to use builtins

- Applications


- Runtimes
- Normalization

- Core area increase by 0.8%

Population Count




Find first 1



Mia Wallace

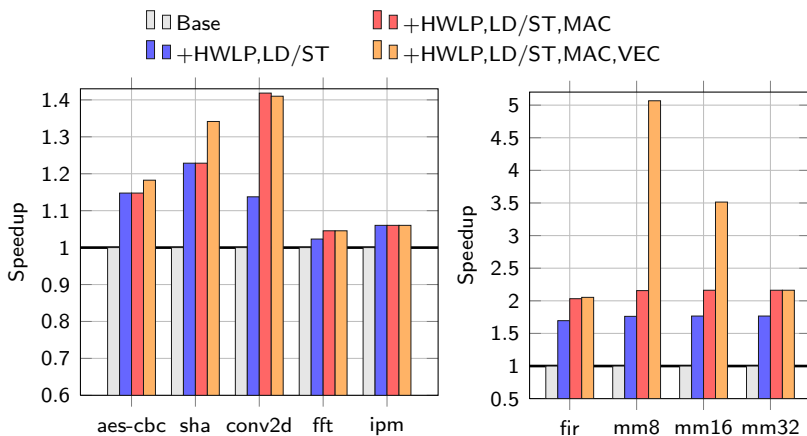
PULP chip with 4 cores in 65 nm



- Max. Frequency: 333 MHz @ 1.08V
- L2 Memory: 256 kB
- TCDM: 64 + 8 kB

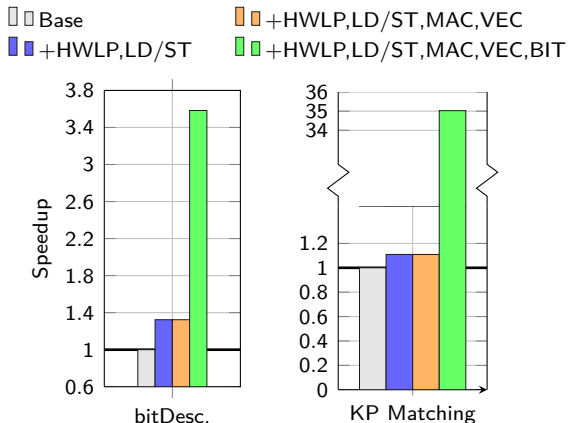
Performance: MAC and Vectorial

Increased up to 5x







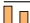

Performance: Bit Count


Increased up to 35x



Area & Timing

Area increase by 25%


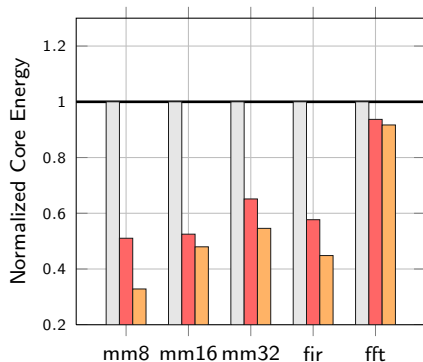

Feature	Area	
 Baseline	35.5 kGE	
 HWLP	3.0 kGE	+8.5%
 Reg. File Add.	3.7 kGE	+10.4%
 New MAC	1.2 kGE	+3.3%
 Vectorial ALU	1.9 kGE	+5.3%
 Bit Count	0.2 kGE	+0.8%
Total	44.5 kGE	



- No increase in critical path

Energy

45% more energy efficient on average



Conclusion

- Vectorial unit and new MAC cost only ~9% core area
Less than 1% cluster area increase
- Up to 5x speedup compared to base OpenRISC on vectorial code
Instructions can be generated automatically by the compiler
- Up to 35x speedup with bit counting instructions
- No increase in critical path
- 45% more energy efficient

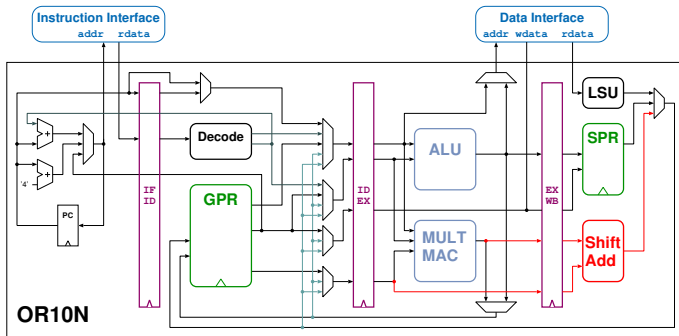
Backup Slides

Outlook

- Look into fractional support

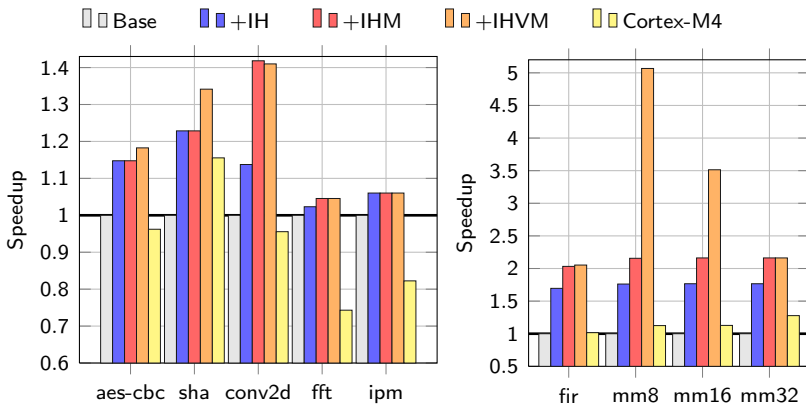
$$rD = (rD \gg 16) + (rA[31 : 16] \cdot rB[31 : 16]) \ll 1$$

- Redesign multiplier



Performance: MAC and Vectorial

Increased up to 5x



Area Multiplier

in 28nm

Design	Features	Area OR10N	Area Mult.	Timing
Baseline		20125	5950	7.57 ns
Vec. Booth Mult.	MAC	18925	5230	7.83 ns
Vec. Booth Mult.	MAC, subword	18850	5451	7.83 ns
Shared Segmentation	MAC	19518	6687	7.77 ns
Shared Segmentation	MAC, subword	18854	5125	7.84 ns
Behavioral	MAC	18900	4202	7.58 ns
Behavioral	MAC, subword	18884	4287	7.64 ns

Vectorial Booth Multiplier

